Chapter 5

Causal Structure Learning II: The PC Algorithm and Sparsity-based Scores

5.1 Faithfulness

Our example in Section 4.2 used the assumption that any conditional independence in $\mathbb{P}_{\mathcal{X}}$ implied a d-separation in causal graph $\mathcal{G}$. The following example shows how this constraint can be violated:

Example 5.1. *Let M be the structural causal model with*

$$
\begin{aligned}
X_1 &= \varepsilon_1 & \varepsilon_1 &\sim \mathcal{N}(0,1) \\
X_2 &= aX_1 + \varepsilon_2 & \varepsilon_2 &\sim \mathcal{N}(0,1) \\
X_3 &= bX_1 + cX_2 + \varepsilon_3 & \varepsilon_3 &\sim \mathcal{N}(0,1)
\end{aligned}
$$

We have $\mathrm{Cov}(X_1, X_3) = ab + c$. Thus, if $c = -ab$, e.g. if $a = b = 1$ and $c = -1$, then we have $\mathrm{Cov}(X_1, X_3) = 0$. In a normal distribution, zero covariance implies independence, and thus if $c = -ab$, we have $X_1 \perp\!\!\!\perp_{\mathbb{P}_{\mathcal{X}}} X_3$ for the entailed distribution $\mathbb{P}_{\mathcal{X}}$. However, X_1 and X_2 are d-connected given $\mathbf{S} = \varnothing$ in the causal graph of M. Thus, $\mathbb{P}_{\mathcal{X}}$ is not faithful to $\mathcal{G}$.

Thus, the correspondence between conditional independences in $\mathbb{P}_{\mathcal{X}}$ and d-separation in $\mathcal{G}$ constitutes a substantive assumption:

Definition 5.1. *We say that $\mathbb{P}_{\mathcal{X}}$ is **faithful** to $\mathcal{G}$ if $\mathcal{I}_{\perp\!\!\!\perp}(\mathbb{P}_{\mathcal{X}}) = \mathcal{I}_{\perp\!\!\!\perp}(\mathcal{G})$.*

Remark 5.1. *A typical justification for the assumption of faithfulness is that, in linear models, it holds for **generic** choices of parameters. Intuitively, this means that the faithfulness assumption is violated only for carefully-selected "degenerate" choices of the parameters. More formally, this means that the set of parameters for which the assumption is violated has a **Lebesgue measure** of zero. In particular, a violation of the faithfulness assumption implies that some polynomial of the edge weights equals zero, e.g. $ab + c = 0$ in Example 5.1. However, the zero set of a non-trivial polynomial has Lebesgue measure zero.*

Despite this justification, there might be statistical issues due to these polynomials being *close* to zero, and thus hard to discern from true zeros. Thus, in causal structure learning, we aim to precisely understand *which* conditional independences in $\mathbb{P}_{\mathcal{X}}$ should imply d-separations in $\mathcal{G}$ - often, only a small subset of conditional independences actually get used by the algorithm. Two particularly important weaker versions of the faithfulness assumptions are adjacency-faithfulness and orientation-faithfulness.

Definition 5.2. *We say that $\mathbb{P}_{\mathcal{X}}$ is **adjacency-faithful** to $\mathcal{G}$ if $X_i \in \mathrm{adj}_{\mathcal{G}}(X_j)$ implies that there is no set $\mathbf{S}$ for which $X_i \perp\!\!\!\perp_{\mathbb{P}_{\mathcal{X}}} X_j \mid \mathbf{S}$.*

Algorithm 1 `PC-Skeleton`

1: **Input:** Conditional independence tester H
2: **Output:** Skeleton $\widehat{\mathcal{G}}$, separator function s
3: Let $\widehat{\mathcal{G}}$ be the complete graph and s be the empty function.
4: Let $d = 0$.
5: **while** $\exists\, X_i - X_j$ in $\widehat{\mathcal{G}}$ such that $|\operatorname{adj}_{\widehat{\mathcal{G}}}(X_i) \setminus \{X_j\}| \geq d$ **do**
6:    **for** $X_i - X_j$ in $\widehat{\mathcal{G}}$ **do**
7:       **if** $\exists\, \mathbf{S} \subseteq \operatorname{adj}_{\mathcal{G}}(X_i) \setminus \{X_j\}$ or $\mathbf{S} \subseteq \operatorname{adj}_{\mathcal{G}}(X_j) \setminus \{X_i\}$ such that $|\mathbf{S}| = d$ and $X_i \perp\!\!\!\perp_H X_j \mid \mathbf{S}$ **then**
8:          Remove $X_i - X_j$ from $\widehat{\mathcal{G}}$
9:          Assign $s(X_i, X_j) = \mathbf{S}$
10:       **end if**
11:    **end for**
12:    Let $d = d + 1$
13: **end while**
14: **return** $\widehat{\mathcal{G}}, s$

Definition 5.3. *We say that $\mathbb{P}_{\mathcal{X}}$ is **orientation-faithful** to $\mathcal{G}$ if, for all $X_i - X_k - X_j$ in $\mathcal{G}$ with $X_i \notin \operatorname{adj}_{\mathcal{G}}(X_j)$, we have:*

- *$X_i \to X_k \leftarrow X_j$ implies that for any $\mathbf{S}$ such that $X_i \perp\!\!\!\perp_{\mathbb{P}_{\mathcal{X}}} X_j \mid \mathbf{S}$, we have $X_k \notin \mathbf{S}$ and*

- *otherwise, for any $\mathbf{S}$ such that $X_i \perp\!\!\!\perp_{\mathbb{P}_{\mathcal{X}}} X_j \mid \mathbf{S}$, we have $X_k \in \mathbf{S}$*

*We say that $\mathbb{P}_{\mathcal{X}}$ is **restricted-faithful** to $\mathcal{G}$ if it is both adjacency-faithful and orientation-faithful to $\mathcal{G}$.*

5.2   The PC algorithm

We are finally ready to introduce the *PC algorithm* (PC stands for "Peter-Clark", the first names of Peter Spirtes and Clark Glymour, the developers of the algorithm). We will introduce the abstraction of a **conditional independence tester**: a function, which given two variables X_i, X_j and a set $\mathbf{S}$, returns a boolean value `True` or `False`. In the *noiseless* (also called *population*) version of the PC algorithm, which assumes direct access to the distribution $\mathbb{P}_{\mathcal{X}}$, we may define a conditional independence tester which simply checks whether $X_i \perp\!\!\!\perp_{\mathbb{P}_{\mathcal{X}}} X_j \mid \mathbf{S}$. In the *sample* version of the PC algorithm, several hypothesis tests exist for testing conditional independence, with different guarantees under different sets of assumptions.

5.2.1   A description of the PC algorithm

The PC algorithm, in Algorithm 3, consists of two phases correspond to the two graphical properties that define Markov equivalence. In its first phase, found in Algorithm 1, the PC algorithm learns the skeleton of $\mathcal{G}$. In its second phase, found in Algorithm 2, the PC algorithm adds orientations to this skeleton.

To learn the skeleton of $\mathcal{G}$, the PC algorithm deletes edges in "rounds". If we assume orientation faithfulness for the conditional independence tester H, then $X_i \perp\!\!\!\perp_H X_j \mid \mathbf{S}$ implies that we can delete the edge $X_i - X_j$. For computational efficiency, we do not want to check all sets $\mathbf{S}$, so each round consider only sets of size d, where $d = 0$ in the first round and increments by one in each following round. Moreover, by Corollary 4.1, it suffices to check only those sets $\mathbf{S}$ which can possibly be parent sets of either X_i or X_j.

The first phase also saves which sets separated each pair X_i, X_j that are non-adjacent in the skeleton. The second phase uses these sets to determine unshielded colliders, orienting some edges. Then, it applies the Meek rules repeatedly to construct the essential graph.

Theorem 5.1. *Suppose that $\mathbb{P}_{\mathcal{X}}$ factorizes according to $\mathcal{G}$, and that $\mathbb{P}_{\mathcal{X}}$ is restricted-faithful to $\mathcal{G}$. Then, in the noiseless case, the PC algorithm returns $\mathcal{E}(\mathcal{G})$.*

Algorithm 2 PC-Orient

1: **Input:** Skeleton $\widehat{\mathcal{G}}$, separator function s
2: **Output:** Essential graph $\widehat{\mathcal{G}}$
3: **for** $X_i - X_k - X_j$ in $\widehat{\mathcal{G}}$ with X_i and X_j non-adjacent **do**
4:     **if** $X_k \notin s(X_i, X_j)$ **then**
5:         Let $X_i \to X_k \leftarrow X_j$ in $\widehat{\mathcal{G}}$
6:     **end if**
7: **end for**
8: Return MPDAG($\widehat{\mathcal{G}}$)

Algorithm 3 PC

1: **Input:** Conditional independence tester H
2: **Output:** Essential graph $\widehat{\mathcal{G}}$
3: Let $\widehat{\mathcal{G}}, s = $ PC-Skeleton(H)
4: Let $\widehat{\mathcal{G}} = $ PC-Orient($\widehat{\mathcal{G}}, s$)
5: Return $\widehat{\mathcal{G}}$

Proof. First, we prove that the skeleton phase of the PC algorithm is correct. Assume that we have not incorrectly removed any edges, and that we remove the edge $X_i - X_j$. Then we have a set $\mathbf{S}$ such that $X_i \perp\!\!\!\perp_H X_j \mid \mathbf{S}$, and thus by adjacency faithfulness, X_i is not adjacent to X_j in $\mathcal{G}$. Thus, we never incorrectly remove any edges.

Conversely, assume that X_i is not adjacent to X_j in $\mathcal{G}$. Then, by Corollary 4.1, we can assume without loss of generality that $X_i \perp\!\!\!\perp_H X_j \mid \mathrm{pa}_{\mathcal{G}}(X_i)$. The algorithm will reach the round where $d = |\mathrm{pa}_{\mathcal{G}}(X_i)|$ since we never incorrectly remove an edge $X_k - X_i$ for $X_k \in \mathrm{pa}_{\mathcal{G}}(X_i)$. In this round, we will consider $\mathbf{S} = \mathrm{pa}_{\mathcal{G}}(X_i) \subseteq \mathrm{adj}_{\mathcal{G}}(X_i) \setminus \{X_j\}$ and see that $X_i \perp\!\!\!\perp_H X_j \mid \mathbf{S}$, thus removing the edge $X_i - X_j$.

Now, we show that the orientation phase is correct. If $X_i \to X_k \leftarrow X_j$, then $X_k \notin s(X_i, X_j)$ by the first condition of orientation faithfulness. Thus, we do not fail to orient any unshielded colliders. Conversely, if $\langle X_i, X_k, X_j \rangle$ is a non-collider, then $X_k \in s(X_i, X_j)$ by the second condition of orientation faithfulness.

The remaining orientations are given by appealing to Theorem 4.3. □

Lemma 5.1. *Let $\mathcal{G}$ be a DAG with maximum degree Δ. Then, in the noiseless case,* PC-Skeleton *terminates with $d \le \Delta + 1$.*

Proof. Let $\Delta_{\mathtt{in}}$ denote the maximum in-degree of $\mathcal{G}$. At the round when $d = \Delta_{\mathtt{in}}$, we have that $\widehat{\mathcal{G}}$ equals the skeleton of $\mathcal{G}$. Thus, $|\mathrm{adj}_{\widehat{\mathcal{G}}}(X_i)| \le \Delta$ for all X_i after this round. Thus when $d = \Delta + 1$, the while loop breaks. □

Finally, we establish the run-time of the PC algorithm. We will use the notion of **query complexity**, which asks how many times the algorithm makes a query to the conditional independence test. This allows us to abstract away what the run-time of the conditional independence tester and obtain a more general result.

Theorem 5.2. *Let $\mathcal{G}$ be a DAG on p nodes with maximum degree Δ. Then, in the noiseless case, the query complexity of the PC algorithm is $\mathcal{O}(p^{\Delta+2})$.*

Proof. At each round, we consider at most $\mathcal{O}(p^2)$ pairs of nodes. For each pair of node, we consider at most $\binom{p}{d'}$ sets $\mathbf{S}$ in the round where $d = d'$. The algorithm terminates after Δ rounds. Thus, for each pair, we consider at most $\sum_{d'=0}^{\Delta} \binom{p}{d'} = \mathcal{O}(p^{\Delta})$ sets. In total, we make at most $\mathcal{O}(p^{\Delta+2})$ queries to H. □

Conditional independence testing. In practice, one has many choices for the conditional independence tester H. If the data is assumed to be linear Gaussian, then the conditional independence $X_i \perp\!\!\!\perp_{\mathbb{P}_{\mathcal{X}}} X_j \mid \mathbf{S}$ is equivalent to the partial correlation $\rho(X_i, X_j \mid \mathbf{S})$ being equal to zero. Given a sample partial correlation

$\widehat{\rho}(X_i, X_j \mid \mathbf{S})$, there is a standard hypothesis test for the null hypothesis $\rho(X_i, X_j \mid \mathbf{S}) = 0$, based on the Fisher z-transformation of $\widehat{\rho}(X_i, X_j \mid \mathbf{S})$. Note that computing the sample partial correlation for $|\mathbf{S}| = d$ takes $\mathcal{O}(d^3)$ time, which is included in a complete runtime analysis of the PC algorithm. In the nonparametric setting, there are several conditional independence tests, see e.g. Zhang et al. (2011) and Shah and Peters (2020). Many of these tests are kernel-based, and thus their computational complexity scales with the number of samples n, creating practical challenges to scaling.

5.3   Sparsity-based Scores   *permutation dictates the orientations*

We now begin our discussion of score-based approaches. We discuss these approaches from the lens of *sparsity*: finding the sparsest graph which fits the observed data. Through this lens, we separate the structure learning problem into two parts: adjacencies and orientations. In the PC algorithm, we found adjacencies first and added orientations based on unshielded colliders and the Meek rules. Here, we take the reverse approach. We start with a permutation π which dictates the orientations of any edges present in the estimated graph, and find the sparsest graph consistent with π which fits the observed data. Then, we search for a permutation π^* for which the corresponding graph is as sparse as possible.   *why sparse as possible is optimal one?*

5.3.1   Minimal I-MAPs

Definition 5.4. *Let $\mathcal{G}$ be a DAG and $\mathbb{P}_\mathcal{X}$ be a distribution. We say that $\mathcal{G}$ is an* **independence map (I-MAP)** *of $\mathbb{P}_\mathcal{X}$ if $\mathcal{I}_{\perp\!\!\!\perp}(\mathcal{G}) \subseteq \mathcal{I}_{\perp\!\!\!\perp}(\mathbb{P}_\mathcal{X})$, i.e., $\mathbb{P}_\mathcal{X}$ is Markov to $\mathcal{G}$. We say that $\mathcal{G}$ is a* **minimal I-MAP** *of $\mathbb{P}_\mathcal{X}$ if no strict subgraph of $\mathcal{G}$ is an I-MAP of $\mathbb{P}_\mathcal{X}$.*

We use the same terminology for graphs, e.g., we say that $\mathcal{G}$ is an I-MAP of $\mathcal{G}'$ if $\mathcal{I}_{\perp\!\!\!\perp}(\mathcal{G}) \subseteq \mathcal{I}_{\perp\!\!\!\perp}(\mathcal{G}')$.

Example 5.2.

what is Graphoid?

Given a permutation, we want a procedure which constructs a minimal I-MAP of $\mathbb{P}_\mathcal{X}$ which is consistent with the permutation. A naïve procedure could enumerate over all 2^p graphs which are consistent with the permutation. However, we can define a much more efficient procedure, which performs one conditional independence query for each pair of nodes $X_i <_\pi X_j$. For this procedure to return a minimal I-MAP, we will introduce the assumption that $\mathbb{P}_\mathcal{X}$ is a graphoid.

Given a permutation π, we will define a procedure which takes in a pair $X_i <_\pi X_j$ and determines whether or not the edge $X_i \to X_j$ is present in the graph associated to that permutation. This procedure requires

Definition 5.5. *A* **semigraphoid** *H over elements $\mathcal{X}$ is a set of disjoint subsets $(\mathbf{A}, \mathbf{B}, \mathbf{S})$, where $(\mathbf{A}, \mathbf{B}, \mathbf{S}) \in H$ is denoted by $\mathbf{A} \perp\!\!\!\perp_H \mathbf{B} \mid \mathbf{S}$, such that the following properties hold:*

- **Symmetry***: $\mathbf{A} \perp\!\!\!\perp_H \mathbf{B} \mid \mathbf{S} \implies \mathbf{B} \perp\!\!\!\perp_H \mathbf{A} \mid \mathbf{S}$*   *indep means d-sep*

- **Decomposition***: $\mathbf{A} \perp\!\!\!\perp_H \mathbf{B} \mid \mathbf{S} \implies \mathbf{A} \perp\!\!\!\perp_H \mathbf{B}' \mid \mathbf{S}$ for $\mathbf{B}' \subseteq \mathbf{B}$*

- **Weak Union***: $\mathbf{A} \perp\!\!\!\perp_H \mathbf{B} \mid \mathbf{S} \implies \mathbf{A} \perp\!\!\!\perp_H \mathbf{B}_1 \mid \mathbf{S} \cup \mathbf{B}_2$ for $\mathbf{B}_1 \subseteq \mathbf{B}$ where $\mathbf{B}_2 = \mathbf{B} \setminus \mathbf{B}'$*   *fix B'<- B1*

- **Contraction***: $\mathbf{A} \perp\!\!\!\perp_H \mathbf{B}_1 \mid \mathbf{S} \cup \mathbf{B}_2$ and $\mathbf{A} \perp\!\!\!\perp_H \mathbf{B}_2 \mid \mathbf{S} \implies \mathbf{A} \perp\!\!\!\perp_H \mathbf{B} \mid \mathbf{S}$ and for $\mathbf{B} = \mathbf{B}_1 \cup \mathbf{B}_2$*

A **graphoid** *is a semigraphoid such that the following property also holds:*

- **Intersection***: $\mathbf{A} \perp\!\!\!\perp_H \mathbf{B}_1 \mid \mathbf{S} \cup \mathbf{B}_2$ and $\mathbf{A} \perp\!\!\!\perp_H \mathbf{B}_2 \mid \mathbf{S} \cup \mathbf{B}_1 \implies \mathbf{A} \perp\!\!\!\perp \mathbf{B} \mid \mathbf{S}$ for $\mathbf{B} = \mathbf{B}_1 \cup \mathbf{B}_2$*

Remark 5.2. *The set $\mathcal{I}_{\perp\!\!\!\perp}(\mathcal{G})$ for either a DAG or an undirected graph $\mathcal{G}$ is a graphoid. The set $\mathcal{I}_{\perp\!\!\!\perp}(\mathbb{P}_\mathcal{X})$ is always a semigraphoid, and it is a graphoid if $\mathbb{P}_\mathcal{X}$ is strictly positive.*

Note that $\mathbb{P}_\mathcal{X}$ can violate the intersection property if it is not positive, as the following example shows.

permutation defines orientation

Figure 5.1: Minimal I-MAPs of a distribution $\mathbb{P}_{\mathcal{X}}$ faithful to $\mathcal{G}_{12345}$. See Example 5.4.

Example 5.3. *Let $B_1 \sim \mathsf{Ber}(0.5)$ and $B_2 \sim \mathsf{Ber}(0.5)$. Let $A = B_1 \oplus B_2$. Then $\mathbb{P}(A = 1) = \mathbb{P}(A = 1 \mid B_1 = 1) = \mathbb{P}(A = 1 \mid B_2 = 1) = 0.5$, i.e., $A \perp\!\!\!\perp_{\mathbb{P}} B_1$. Symmetrically, $A \perp\!\!\!\perp_{\mathbb{P}} B_2$. However, $A \not\perp\!\!\!\perp_{\mathbb{P}} (B_1, B_2)$. Note that $\mathbb{P}(A = 1, B_1 = 0, B_2 = 0) = 0$, i.e., $\mathbb{P}$ is not positive.*

Proposition 5.1. *Let $\mathbb{P}_{\mathcal{X}}$ be a distribution and π a permutation. Let $\mathcal{G}_{\pi}$ be the DAG with edges $X_i \to X_j$ for $X_i <_{\pi} X_j$ if*

$$X_i \not\perp\!\!\!\perp X_j \mid \mathrm{pre}_{\pi}(X_j) \setminus \{X_i\}. \quad \text{\color{red}{if not indep we give edge}}$$

If $\mathbb{P}_{\mathcal{X}}$ is a graphoid, then $\mathcal{G}_{\pi}$ is a minimal I-MAP of $\mathbb{P}_{\mathcal{X}}$.

Proof. It suffices to show that $\mathbb{P}_{\mathcal{X}}$ factorizes according to $\mathcal{G}_{\pi}$. For this, it is sufficient to show that for all X_j,

$$X_j \perp\!\!\!\perp_{\mathbb{P}_{\mathcal{X}}} \mathrm{pre}_{\pi}(X_j) \setminus \mathrm{pa}_{\mathcal{G}}(X_j) \mid \mathrm{pa}_{\mathcal{G}_{\pi}}(X_j) \tag{5.1}$$

Let $X_u, X_v \in \mathrm{pre}_{\pi}(X_j) \setminus \mathrm{pa}_{\mathcal{G}}(X_j)$. By the intersection property with $\mathbf{S} = \mathrm{pre}_{\sigma}(X_j) \setminus \{X_u, X_v\}$, $\mathbf{B}_1 = \{X_u\}$, and $\mathbf{B}_2 = \{X_v\}$, we have

$$X_j \perp\!\!\!\perp_{\mathbb{P}_{\mathcal{X}}} X_u, X_v \mid \mathrm{pre}_{\sigma}(X_j) \setminus \{X_u, X_v\}$$

Equation (5.1) follows by induction. $\qquad\square$

Remark 5.3. *In practice, constructing $\mathcal{G}_{\pi}$ according to the definition in Proposition 5.1 is not ideal: in a graph with p variables, the conditioning sets used in conditional independence tests become of size $\mathcal{O}(p)$ for nodes later in the permutation. More practical constructions can be made which take advantage of sparsity in the underlying graph.*

Example 5.4. *Suppose that $\mathbb{P}_{\mathcal{X}}$ is faithful to $\mathcal{G}_{12345}$ in Figure 5.1. Then $\mathcal{G}_{23145}$ has the extra adjacency $X_2 \to X_3$, since $X_2 \not\perp\!\!\!\perp X_3$. $\mathcal{G}_{24135}$ has two extra adjacencies, $X_2 \to X_3$, since $X_2 \not\perp\!\!\!\perp X_3 \mid X_4$, and $X_4 \to X_1$, since $X_4 \not\perp\!\!\!\perp X_1 \mid X_2$.*

Proposition 5.2. *Let $\mathbb{P}_{\mathcal{X}}$ be Markov to $\mathcal{G}$, and let π be a <mark>topological</mark> order of $\mathcal{G}$. Then $\mathcal{G}_{\pi} \subseteq \mathcal{G}$.*

Proof. Let $X_i <_{\pi} X_j$ with $X_i \notin \mathrm{pa}_{\mathcal{G}}(X_j)$. Let $\mathbf{V} = \overline{\mathrm{pre}}_{\pi}(X_j)$, which is ancestral by definition of a topological order. Then $\mathbb{P}_{\mathcal{X}}(\mathbf{V})$ factorizes according to $\mathcal{G}[\mathbf{V}]$. All paths from X_i to X_j are blocked in $\mathcal{G}[\mathbf{V}]$ by $\mathrm{pre}_{\pi}(X_j) \setminus \{X_i\}$, since this set includes $\mathrm{pa}_{\mathcal{G}}(X_i)$. Thus, $X_i \perp\!\!\!\perp_{\mathbb{P}_{\mathcal{X}}} X_j \mid \mathrm{pre}_{\pi}(X_j)$, hence G_{π} does not contain the edge $X_i \to X_j$. $\qquad\square$

5.3.2   Sparsest Permutation

Lemma 5.2. *Let $\mathbb{P}_\mathcal{X}$ be Markov to a DAG $\mathcal{G}$. Let $\widehat{\pi} \in \arg\min_\pi |\mathcal{G}_\pi|$, where $|\mathcal{G}_\pi|$ denotes the number of edges in $\mathcal{G}_\pi$. Then*

(a) *$|\mathcal{G}_{\widehat{\pi}}| \leq |\mathcal{G}|$*

(b) *If $\mathbb{P}_\mathcal{X}$ is adjacency-faithful to $\mathcal{G}$, then $\mathrm{skel}(\mathcal{G}_{\widehat{\pi}}) = \mathrm{skel}(\mathcal{G})$.*

(c) *If $\mathbb{P}_\mathcal{X}$ is restricted-faithful to $\mathcal{G}$, then $\mathcal{G}_{\widehat{\pi}} \in \mathcal{M}(\mathcal{G})$.*

Proof.
(a) By Proposition 5.2.

(b) By (a), $\mathrm{skel}(\mathcal{G}_{\widehat{\pi}})$ cannot have more edges than $\mathcal{G}$, so it suffices to show that $\mathrm{skel}(\mathcal{G}) \subseteq \mathrm{skel}(\mathcal{G}_{\widehat{\pi}})$. This is guaranteed by adjacency faithfulness and the definition of $\mathcal{G}_{\widehat{\pi}}$.

(c) By (b), we have $\mathrm{skel}(\mathcal{G}_{\widehat{\pi}}) = \mathrm{skel}(\mathcal{G})$. By Theorem 4.2, it suffices to show that $\mathcal{G}_{\widehat{\pi}}$ and $\mathcal{G}$ also have the same unshielded colliders.

Let $X_i \to X_k \leftarrow X_j$ be unshielded in $\mathcal{G}$. Consider π such that $X_k <_\pi X_j$. If $X_i <_\pi X_j$, then $X_k \in \mathbf{S}$ for $\mathbf{S} = \mathrm{pre}_\pi(X_j) \setminus \{X_i\}$. By orientation-faithfulness, we have that $X_i \not\perp\!\!\!\perp_{\mathbb{P}_\mathcal{X}} X_j \mid \mathbf{S}$, so $X_i \to X_j$ in $\mathcal{G}_\pi$. Symmetrically, if $X_j <_\pi X_i$, we have $X_j \to X_i$ in $\mathcal{G}_\pi$. Thus, $\mathrm{skel}(\mathcal{G}_\pi) \neq \mathrm{skel}(\mathcal{G})$. The case where $X_k <_\pi X_i$ is symmetric. Thus, if $\mathrm{skel}(\mathcal{G}_{\widehat{\pi}}) = \mathrm{skel}(\mathcal{G})$, then we must have $X_i <_{\widehat{\pi}} X_k$ and $X_j <_\pi X_k$, i.e., any unshielded collider in $\mathcal{G}$ is an unshielded collider in $\mathcal{G}_{\widehat{\pi}}$.

Conversely, suppose $X_i - X_k - X_j$ is a non-collider in $\mathcal{G}$, with X_i and X_j non-adjacent. Consider π such that $X_i <_\pi X_k$ and $X_j <_\pi X_k$. If $X_i <_\pi X_j$, then $X_k \notin \mathbf{S}$ for $\mathbf{S} = \mathrm{pre}_\pi(X_j) \setminus \{X_i\}$, so by orientation faithfulness, $X_i \to X_j$ in $\mathcal{G}_\pi$. Thus, $\mathrm{skel}(\mathcal{G}_\pi) \neq \mathrm{skel}(\mathcal{G})$. The case where $X_j <_\pi X_i$ is symmetric. Thus, if $\mathrm{skel}(\mathcal{G}_{\widehat{\pi}}) = \mathrm{skel}(\mathcal{G})$, then we must have at least one of $X_k <_\pi X_i$ or $X_k <_\pi X_j$, i.e., any unshielded non-collider in $\mathcal{G}$ is an unshielded non-collider in $\mathcal{G}_{\widehat{\pi}}$. □

5.4   Greedy Search Algorithms

Definition 5.6. *Let $\mathcal{G}$ be a DAG. We call an edge X_i to X_j in $\mathcal{G}$ a **covered edge** if $\overline{\mathrm{pa}}_\mathcal{G}(X_i) = \mathrm{pa}_\mathcal{G}(X_j)$, equivalently, $\mathrm{pa}_\mathcal{G}(X_j) \setminus \mathrm{pa}_\mathcal{G}(X_i) = \{X_i\}$.*

Lemma 5.3. *Let $\mathcal{G}$ be a DAG and let $X_i \to X_j$ be a covered edge. Let $\mathcal{G}'$ be the graph obtained by reversing the edge $X_i \to X_j$ in $\mathcal{G}$, i.e., in $\mathcal{G}'$, we have $X_j \to X_i$. Then $\mathcal{G}'$ is Markov equivalent to $\mathcal{G}$.*

Proof. Note that, by acyclicity, there are no directed paths from X_i to X_j except for the edge $X_i \to X_j$. Thus, no new cycles are introduced by reversing $X_i \to X_j$, i.e., $\mathcal{G}'$ is a DAG.

The skeletons of $\mathcal{G}$ and $\mathcal{G}'$ are equal. So are their unshielded colliders: the only colliders in $\mathcal{G}'$ that are not in $\mathcal{G}$ are of the form $X_k \to X_i \leftarrow X_j$, but these are all shielded since by the assumption that the edge is covered. Symmetrically, the only colliders in $\mathcal{G}$ that are not in $\mathcal{G}$ are shielded. Thus, by Theorem 4.2, $\mathcal{G}$ and $\mathcal{G}'$ are Markov equivalent. □

In fact, it is well-known that for *any* $\mathcal{G}'$ that is Markov equivalent to $\mathcal{G}$, there exists a sequence of covered edge reversals from $\mathcal{G}$ to $\mathcal{G}'$. This is a special case of the upcoming result, which is one of the major results in score-based causal structure learning.

Definition 5.7. *Let $\mathcal{G}$ be a DAG and $\mathcal{H}$ be an I-MAP of $\mathcal{G}$, i.e., $\mathcal{I}_\perp(\mathcal{H}) \subseteq \mathcal{I}_\perp(\mathcal{G})$. A **Chickering sequence** from $\mathcal{G}$ to $\mathcal{H}$ is a sequence of DAGs $\mathcal{G}_0, \mathcal{G}_1, \ldots, \mathcal{G}_M$, with $\mathcal{G}_0 = \mathcal{G}$ and $\mathcal{H} = \mathcal{G}_M$, such that $\mathcal{G}_m$ is obtained from $\mathcal{G}_{m-1}$ by either an edge addition or a covered edge reversal.*

Theorem 5.3 (Chickering (2002))**.** *Let $\mathcal{H}$ be an I-MAP of $\mathcal{G}$. Then there is a Chickering sequence from $\mathcal{G}$ to $\mathcal{H}$.*

Figure 5.2: A Chickering sequence from $\mathcal{G}$ to $\mathcal{H}$. Blue edges denoted covered edges.

Remark 5.4. *Chickering (2002) offers a constructive proof of Theorem 5.3. Given a pair $\mathcal{G}_m$ and $\mathcal{H}$, the author introduces an operation which either reverses a covered edge in $\mathcal{G}_m$ or adds an edge to $\mathcal{G}_m$, obtaining a new graph $\mathcal{G}_{m+1}$ such that $\mathcal{H}$ is still an I-MAP of $\mathcal{G}_{m+1}$. The proof defines a distance $d(\mathcal{G}, \mathcal{H})$ between graphs, and shows that $d(\mathcal{G}_{m+1}, \mathcal{H}) < d(\mathcal{G}_m, \mathcal{H})$. Thus, repeated application of the operation must terminate when $d(\mathcal{G}_M, \mathcal{H}) = 0$, i.e., $\mathcal{G}_M = \mathcal{H}$. Since this proof would take a lecture of its own to go through, we will not cover it in this course.*

Example 5.5. *An example Chickering sequence can be found in Figure 5.2.*

Corollary 5.1. *Let $\mathcal{H}$ be an I-MAP of $\mathcal{G}$ such that $\mathcal{I}_{\perp\!\!\!\perp}(\mathcal{H}) \neq \mathcal{I}_{\perp\!\!\!\perp}(\mathcal{G})$. Then there exists $\mathcal{H}' \in \mathcal{M}(\mathcal{H})$ such that $\mathcal{H}'$ is not a minimal I-MAP of $\mathcal{G}$.*

Proof. By Theorem 5.3, there exists a Chickering sequence from $\mathcal{G}$ to $\mathcal{H}$. Since $\mathcal{I}_{\perp\!\!\!\perp}(\mathcal{H}) \neq \mathcal{I}_{\perp\!\!\!\perp}(\mathcal{G})$, we must have $|\mathcal{H}| > |\mathcal{G}|$. Let $\mathcal{G}_m$ be the last graph in this sequence such that $|\mathcal{G}_m| = |\mathcal{H}| - 1$. Then $\mathcal{H}' = \mathcal{G}_{m+1}$ is not a minimal I-MAP. □

Two greedy search procedures exploit this corollary. First, the **Greedy Sparsest Permutation** (GSP)

algorithm heuristically picks an initial permutation π_0, and uses conditional independence tests to estimate $\widehat{\mathcal{G}}_{\pi_0}$. Then, GSP uses a depth-first search (via covered edge reversals) over the Markov equivalence class of $\widehat{\mathcal{G}}_{\pi_0}$ to find some graph for which an edge can be deleted. This can be seen as a "backwards" traversal of the Chickering sequence. Further, the use of covered edge reversals allows for very few conditional independence tests to be performed at each step, i.e., the minimal I-MAP $\mathcal{G}_\pi$ does not have to be re-computed at each iteration. GSP can be seen as a "hybrid" method, using conditional independence tests for each fixed permutation, with a score-based search over permutations which looks to minimize the number of edges.

The **Greedy Equivalence Search** (GES) algorithm has both a forward and a backward phase. The forward phase, which starts from the empty graph and adds edges, is designed to end at an I-MAP of the $\mathbb{P}_{\mathcal{X}}$, while the backward phase follows a similar path "down" the Chickering sequence. Instead of depth-first search within a Markov equivalence class, GES operates over the search essential graphs, and thus involves somewhat complex rules for adding and deleting edges in the forward and backward phases. Finally, we note that GES is usually run with a score function which takes the form of either a penalized maximum likelihood or a marginal likelihood (where some prior is defined over graphs and parameters).

5.5   Additional Remarks

- **Causal structure learning with uncertainty**: The algorithms that we discussed return "point estimates", i.e., they output a single estimated graph $\widehat{\mathcal{G}}$. However, in practice, one may want to quantify the uncertainty over this estimated graph. A frequentist approach is to use a *DAG-bootstrap* (Friedman et al., 1999), which re-runs one of these algorithms with bootstrapped samples of the data (i.e., random sampling with replacement). Alternatively, one may take a Bayesian approach, defining priors $\mathbb{P}(\mathcal{G})$ and $\mathbb{P}(\Theta \mid \mathcal{G})$ over causal graphs and their parameters. You can read more about some of these approaches in Section 4.2 of my review (Squires and Uhler, 2022).

- **Causal structure learning with interventions**: Many causal structure learning algorithms can be re-purposed for learning with interventional data. As indicated by our definition of an interventional graph in Chapter 1, this can be done by adding intervention variables as nodes Mooij et al. (2020). Such an approach also allows one to learn the *targets* of an intervention, i.e. the targets do not need to be known beforehand. The main criterion for a structure learning algorithm to be re-purposed to handle interventional data is that it can incorporate *background knowledge*. In particular, the algorithm must be able to incorporate the facts that intervention variables are upstream of the "normal" variables, and that intervention variables (if there are multiple) have a deterministic relation (if $\zeta^I = 1$ for some sample, then $\zeta^{I'} = 0$ for any other intervention I'). An example of such an adaptation of the GSP algorithm is in Squires et al. (2020).

- **Causal structure learning with latent variables**: Here, we have covered causal structure learning for DAG models, i.e., in the setting of a Markovian structural causal model. There are many methods extending beyond DAGs, e.g. to handle latent variables. There are several options for graphical models which capture latent variables, including ADMGs (discussed in Chapter 1) or **maximal ancestral graphs** (MAGs). For example, the **Fast Causal Inference** (FCI) algorithm and its variants are analogues of the PC algorithm for learning MAGs. One barrier to proving the consistency of greedy search algorithms is that there is no proof that a Chickering sequence always exists between a MAG $\mathcal{G}$ and an I-MAP $\mathcal{H}$ of that MAG. For example, **Greedy Sparsest Poset** (GSPo) algorithm is a variant of the GSP algorithm for learning MAGs, but its consistency remains unproven.

Bibliography

Chickering, D. M. (2002). Optimal structure identification with greedy search. *Journal of machine learning research*, 3(Nov):507–554.

Friedman, N., Goldszmidt, M., and Wyner, A. (1999). Data analysis with bayesian networks: a bootstrap

approach. In *Proceedings of the Fifteenth conference on Uncertainty in artificial intelligence*, pages 196–205.

Mooij, J. M., Magliacane, S., and Claassen, T. (2020). Joint causal inference from multiple contexts.

Shah, R. D. and Peters, J. (2020). The hardness of conditional independence testing and the generalised covariance measure. *The Annals of Statistics*, 48(3):1514–1538.

Squires, C. and Uhler, C. (2022). Causal structure learning: a combinatorial perspective. *arXiv preprint arXiv:2206.01152*.

Squires, C., Wang, Y., and Uhler, C. (2020). Permutation-based causal structure learning with unknown intervention targets. In *Conference on Uncertainty in Artificial Intelligence*, pages 1039–1048. PMLR.

Zhang, K., Peters, J., Janzing, D., and Schölkopf, B. (2011). Kernel-based conditional independence test and application in causal discovery. In *Proceedings of the Twenty-Seventh Conference on Uncertainty in Artificial Intelligence*, pages 804–813.